ChatGPT Response and User feedback 1:

1. Product scope and users

1. Is v1 **mobile-only**, **web-only**, or **both mobile and web**?
   Your doc mentions both phone and web, but later says mobile native apps are out of scope for v1, while also requiring iOS/Android deployment. That needs to be resolved.

Answer: Let us develop the product for both mobile and web.

2. Who are the actual user roles in v1?
   - Therapist
   - Patient
   - Practice admin
   - Supervisor/reviewer
   - Billing staff

Answer: Therapist is the actual user of the product. They register with the app once and the backend will use the registration details to authenticate the user's login, map the Azure blob store folder structure.

Also before uploading the audio file after every therapy session, the therapist would enter the patient name and their appointment date and time. This will be used for mapping the audio file and the final output all into a folder that will uniquely be mapped by therapist, patient, and the session date/time.

3. Will **patients** do anything in the app besides registration, or is the app primarily for therapists?

Answer: Yes, the app is primarily for therapists

4. Is this for **independent therapists**, **small practices**, or **larger multi-clinic organizations**?

Answer: This is for both independent and small practices

2. Compliance and legal

5. Are you planning to build this as a **HIPAA-ready product from day one**, or do you want a **prototype / MVP first** without production PHI?

Answer: Yes, the product will be HIPAA-ready product from day one. The audio file will be recorded by the therapist after the session and would not involve the conversation with the patient.

6. Do you already have access to:
    - an **Azure subscription**
    - a **BAA with Microsoft/Azure**
    - a **BAA with the LLM provider** you intend to use?

Answer: yes

7. When you said "use a general purpose LLM like ChatGPT 4-mini," do you specifically mean:

    - **Azure OpenAI**, or
    - direct OpenAI API?

Answer: Azure OpenAI

8. For real PHI, this is a big architectural decision. Your requirements say AI providers must be BAA-compliant and not use data for training.

Answer: yes

9. Which country/state will the app operate in first?
   That affects data residency, consent language, and recording rules.

Answer: USA

10. Will the therapist always obtain **documented patient consent** for session recording?

Answer: The actual session would not be recorded but the therapist will record the details after the session is over. There would be any references to the patient in the database and the only details that will be maintained related to the session will the session date and time which the therapist can use to link to the patient.

---

3. Audio capture and transcription

10. Will the uploaded audio be:

- therapist-only voice summary after session,

Answer: yes.

- full session recording with both therapist and patient,

Answer: No

- or both?Answer: No

The doc says “clinical audio recording from just the clinician,” but some sections also mention diarization and full-session processing. Those imply different pipelines.

11. Do you really need **speaker diarization** in v1?
    If therapist records a post-session verbal summary alone, diarization may be unnecessary.

    Answer: Yes

12. What is the expected average audio length:

- 5–10 min therapist dictation
- 45–60 min session recording
- up to 120 min regularly?

Answer: up to 60 mins from session

13. Should therapists be able to:

- record directly in-app,
- upload from Files/Voice Memos,
- or both?

Answer: directly record in the app. But for the mvp they can use another voice recorder app in their phone and upload it via our app.

14. Do you need multilingual support, accents, or only English for v1?

Answer: supports only English in MVP. Later will be expanded to other languages

---

4. SOAP note generation workflow

15. Should the system generate:

- only a SOAP note,
- SOAP + progress note,
- SOAP + intervention suggestions,
- SOAP + downloadable DOCX/PDF?

Answer: SOAP and downloadable DOCX

16. Should the generated note be:

- **strictly extractive** from transcript only,
- or allowed to make **light clinical summarization/inference**?

Answer: The generated note can have a light clinical summarization inference.

This matters because your examples include clinically polished synthesis, but your requirements also say never alter meaning and flag uncertain content.

17. Do you want the Assessment and Plan sections to include:

- CBT
- DBT
- Choice Reality Therapy
- Gottman
- Solution-Focused
- custom therapist-selected frameworks?

Answer: yes

18. Should the therapist choose the therapy modalities **before generation**, or should the AI infer likely interventions?

Answer: For MVP, let the AI infer likely interventions

19. Do you want the app to pull in:

- previous session notes,
- treatment plan,
- diagnosis,
- medications,
- risk history,
- homework/completed interventions?

Answer: not for MVP release

Your problem statement says it should connect prior notes and treatment plan, which makes this more than a simple audio-to-note app.

20. Should the generated output highlight:

- missing evidence,
- risky assumptions,
- hallucination warnings,
- PHI mentions,
- confidence per section?

Answer: yes, please

---

5. Review, editing, and document lifecycle

21. Where should therapist review happen?

- inside the app
- exported Word document only
- both

22. Do you want a **rich text editor** for note review in the app?
23. Should edits made by therapist be saved as:

- final version only
- version history / audit trail of AI output vs human edits

24. Do therapists need templates customized by practice or by clinician?
25. Should the final note be:

- downloadable only,
- stored in system,
- emailed,
- pushed to EHR later?

Answer: The therapist will have a dashboard in the app, where they can see the output of the backend and can nudge or request a revision

---

6. Identity, tenancy, and data model

26. Is this a **single-practice** app first, or a **multi-tenant SaaS** from the beginning?

27. Can one therapist belong to multiple practices?
28. Can one patient have multiple therapists?
29. Do you need family/couples therapy later, or should data model stay strictly one therapist ↔ one patient for v1?
30. What minimum patient profile fields are required at registration?

- legal name
- DOB
- insurance info
- diagnosis
- emergency contact
- consent flags

Answer: For MVP let us focus on single-practice app first

No patient details are collected in the app.

Couple/families therapies will be implemented later.

---

7. Backend and deployment

31. For Python backend on Azure, are you leaning toward:

- **FastAPI**
- **Azure Functions**
- or a hybrid approach?

Answer: I already developed a backend proof-of-concept using FastAPI

32. Do you want the backend to be primarily:

- synchronous API calls for short jobs,
- or async job pipeline with queue/orchestration for transcription + LLM + document generation?

Answer: async job pipeline

33. Are you expecting **50+ concurrent sessions in production from day one**, or is that a later scaling goal?

Answer: scaling would be a later consideration.

34. Is there a target budget for MVP cloud spend per month?

Answer: $500 initial monthly budget.

35. Do you need separate environments:

- dev
- test
- staging
- prod

Answer: Dev only for POC / MVP

---

8. Security

36. Should therapists authenticate with:

- email/password + MFA
- Microsoft/Google SSO
- practice-managed identity
- Entra External ID / B2C type login

Answer: Entra External ID for MVP. later can be integrated with MS / Google SSO

37. Do you need role-based access such that therapists can only see their own patients, while admins can see practice-wide data? Your requirements suggest yes, but I'd want that confirmed.

Answer: Yes, therapists can see only their data and admins can see all data.

38. Do you need full audit logs for:

- upload
- transcription
- note generation
- edits
- downloads
- Deletes

Answer: yes

39. What is the retention policy for:

- raw audio
- raw transcript
- cleaned transcript
- generated SOAP note
- final document

Answer: 1 year retention policy

---

9. Business and go-to-market

40. What is the intended pricing model for v1?

- per therapist / month
- per session processed
- practice subscription
- Hybrid

Answer: per session processed or a monthly subscription

41. Is the real goal:

- build a production startup product,
- build an MVP for pilot users,
- or create a demo/prototype for fundraising?

Answer: build a MVP for pilot users and for fundraising

42. How many initial pilot users do you expect in the first 3 months?

Answer: 2

---

10. Output you want from me next

43. What do you want me to produce after you answer?

- solution architecture
- Azure service design
- API design
- Flutter + Python folder structure
- MVP roadmap
- security/compliance checklist
- cost estimate
- pitch-deck style startup recommendation
- all of the above in one package

Start with the solution architecture, and ask me after that.

ChatGPT Response 2:

Perfect — with your answers, here's the **recommended MVP solution architecture** for the therapy SOAP-note product.

1. Recommended architecture at a glance

Build the MVP as a **Flutter client** for iPhone + web, backed by a **Python FastAPI API** on Azure, with an **asynchronous processing pipeline** for audio upload → transcription → cleanup → SOAP generation → therapist review. For your use case, Azure gives you the right building blocks for HIPAA-oriented design: **Microsoft Entra External ID** for customer sign-in, **Blob Storage** for audio and generated files, **Azure AI Speech batch transcription** for async speech-to-text, and **Azure OpenAI** for note generation. Azure documents also state Azure OpenAI customer prompts/completions aren't used to retrain models, and Azure Speech supports async batch transcription workflows. (Microsoft Learn)

High-level flow

1. Therapist signs in from Flutter app.
2. Therapist enters:
    - patient alias/name
    - appointment date/time
    - optional note context
3. Therapist uploads audio file.
4. FastAPI creates a **session job** and stores metadata.
5. Audio lands in **Azure Blob Storage**.

6. Backend submits transcription job to **Azure AI Speech**.
7. When transcript is ready, backend:
    - cleans transcript
    - runs PHI/quality checks
    - calls **Azure OpenAI**
    - generates:
        - SOAP note
        - review flags
        - confidence/assumption warnings
8. Backend stores:
    - raw transcript
    - cleaned transcript
    - generated SOAP JSON
    - DOCX output
    - audit trail
9. Therapist opens dashboard and:
    - reviews draft
    - nudges/regenerates
    - edits/finalizes
10. Final note stays stored for retention and can be downloaded.

---

2. Proposed Azure architecture

A. Client layer

Flutter app

- iPhone app
- responsive web app
- same codebase for MVP

Core therapist screens

- sign up / sign in
- dashboard
- new session
- upload audio
- processing status
- SOAP review / revise
- download DOCX
- audit/history

B. Identity and access

Microsoft Entra External ID

- therapist authentication
- practice admin authentication later
- MFA-capable customer identity flow
- app issues tokens used by FastAPI backend (Microsoft Learn)

Authorization model

- **therapist**: only sees their own sessions and documents
- **practice admin**: sees all therapist sessions within that practice
- single-practice MVP, but design tenant-aware IDs from day one

---

C. API and application layer

FastAPI backend

- keep your existing proof of concept
- containerize it
- deploy on **Azure Container Apps** for MVP

Why this is the best fit:

- works well with Python/FastAPI
- easier than full AKS for MVP
- better for async background workers than a purely synchronous API
- supports container-based web/API deployment directly for FastAPI apps (Microsoft Learn)

Recommended FastAPI components

- **API service**
  - auth validation
  - session CRUD
  - upload orchestration
  - dashboard APIs
- **worker service**
  - transcription job execution
  - transcript cleanup
  - LLM orchestration

    - DOCX generation
    - retries / error handling

D. Storage layer

Azure Blob Storage
Use separate containers or virtual paths for:

- `audio-raw/`
- `transcripts-raw/`
- `transcripts-clean/`
- `soap-json/`
- `soap-docx/`
- `audit-exports/`

Azure Blob supports lifecycle management and immutability options, which fit your retention and compliance needs. (Microsoft Learn)

Recommended folder/key pattern
Use IDs, not plain patient names, in storage paths.

Example:

```
/practice/{practice_id}/therapist/{therapist_id}/patient/{patient_ref}/session/{session_id}/
  audio/raw.m4a
  transcript/raw.json
  transcript/clean.txt
  soap/draft_v1.json
  soap/draft_v1.docx
  soap/final_v1.docx
  audit/events.json
```

Important recommendation
Even though the therapist enters “patient name” in the UI, do **not** use raw patient name in blob path names if you can avoid it. Use:

- `patient_ref`
- `session_id`
- encrypted or pseudonymous mapping in database

That reduces accidental PHI exposure in logs, URLs, and admin tooling.

E. Database layer

Use **Azure Database for PostgreSQL** for MVP.

Store:

- therapists
- practices
- patient references
- sessions
- transcription jobs
- LLM generations
- review actions
- audit events
- document metadata

Why PostgreSQL

- relational structure fits sessions, users, and audit trails
- easier reporting than NoSQL for this app
- good fit with FastAPI
- clean path to multi-tenant SaaS later

F. Async processing layer

Since you want an async pipeline, use **Azure Service Bus** between API and workers. Microsoft's docs position Service Bus as Azure's richer queueing/messaging option compared with basic storage queues. (Microsoft Learn)

Recommended queues

- `session-uploaded`
- `transcription-requested`
- `transcription-completed`
- `soap-generation-requested`
- `docx-generation-requested`
- `job-failed`

Why Service Bus instead of basic storage queues

- better retry/dead-letter workflow

- stronger enterprise messaging semantics
- easier future expansion when you add more workflow stages

G. Transcription layer

Use **Azure AI Speech batch transcription** for MVP because your workflow is naturally asynchronous: user uploads audio, backend processes it later, results are retrieved when ready. Azure's Speech docs explicitly support batch transcription and async result retrieval. (Microsoft Learn)

Recommendation
For MVP, treat recordings as **single-speaker therapist dictation**.

That means:

- no diarization needed
- lower complexity
- lower cost
- less error surface

I know you answered "yes" for diarization, but I'd strongly recommend **not implementing diarization in MVP**. It adds work without meaningful value for therapist-only post-session audio.

H. LLM orchestration layer

Use **Azure OpenAI** for:

- transcript cleanup prompt
- structured SOAP generation
- intervention inference
- "warning flags" generation
- confidence/risk annotation

Azure's docs say Azure OpenAI doesn't use customer data to retrain models, and Microsoft emphasizes prompt engineering plus downstream validation because outputs can still be imperfect. (Microsoft Learn)

Recommended generation pattern
Use **multi-step prompting**, not one giant prompt:

1. **Cleanup step**
    - normalize transcript

        - remove filler
        - preserve meaning
        - mark uncertainty
    2. **Structured extraction step**
        - extract facts into JSON:
            - symptoms
            - behaviors
            - interventions
            - risks
            - homework
            - follow-up
    3. **SOAP generation step**
        - generate SOAP from extracted JSON
        - produce section-level confidence
        - flag assumptions
    4. **DOCX formatting step**
        - populate a controlled template

This pattern is safer than “audio transcript directly to polished note.”

---

3. Recommended logical architecture diagram

```
Flutter iPhone App / Flutter Web
        |
        v
Microsoft Entra External ID
        |
        v
FastAPI API (Azure Container Apps)
        |
        +--> Azure Database for PostgreSQL
        |
        +--> Azure Blob Storage
        |
        +--> Azure Service Bus
                |
                v
         Worker Service (Azure Container Apps / background worker)
                |
                +--> Azure AI Speech (batch transcription)
                |
                +--> Transcript cleanup service
```

```
|
+--> Azure OpenAI
|
+--> DOCX generator
|
+--> Blob Storage + PostgreSQL audit updates
```

---

4. Session data model for MVP

Core entities

Practice

- `practice_id`
- `name`
- `status`

Therapist

- `therapist_id`
- `practice_id`
- `external_identity_id`
- `email`
- `display_name`
- `license_type`
- `created_at`

PatientRef

- `patient_ref_id`
- `practice_id`
- `therapist_id`
- `display_label`
- optional pseudonymous code
- no full patient profile for MVP

Session

- `session_id`
- `practice_id`

- `therapist_id`
- `patient_ref_id`
- `appointment_start_at`
- `appointment_end_at` optional
- `audio_status`
- `transcription_status`
- `soap_status`
- `final_status`
- `created_at`

Artifact

- `artifact_id`
- `session_id`
- `type` (`audio_raw`, `transcript_raw`, `transcript_clean`, `soap_json`, `soap_docx`)
- `blob_uri`
- `version`
- `created_at`

AuditEvent

- `event_id`
- `session_id`
- `actor_type` (`therapist`, `system`, `admin`)
- `event_type`
- `timestamp`
- `metadata_json`

---

5. Recommended therapist workflow

Step 1: register / login

- therapist creates account
- backend maps therapist to practice
- backend provisions logical storage ownership

Step 2: create session

Therapist enters:

- patient alias/name
- appointment date/time
- optional free-text tags

Step 3: upload audio

- app uploads audio to backend
- backend writes to blob
- backend returns session job ID

Step 4: process

Background worker:

- validates audio format
- submits speech batch job
- stores transcript
- runs cleanup
- runs SOAP generation
- builds DOCX
- updates session status

Step 5: therapist review

Dashboard shows:

- processing state
- extracted transcript summary
- generated SOAP note
- confidence/assumption flags
- nudge buttons:
    - “Make more concise”
    - “Strengthen plan section”
    - “Reduce assumptions”
    - “Regenerate with CBT framing”

Step 6: finalize

- therapist edits and approves
- final DOCX stored
- audit trail updated

---

6. Security architecture for MVP

Because you want HIPAA-ready design from day one, the minimum architecture should include:

Identity

- Entra External ID for auth
- MFA enabled for therapist accounts
- short-lived access tokens

Encryption

- TLS in transit
- Azure-managed encryption at rest at minimum
- prefer customer-managed keys later for stronger enterprise posture

Access control

- therapist-level row filtering
- practice admin scoped separately
- no cross-therapist access by default

Secrets

- store all secrets in **Azure Key Vault**
- never hardcode:
    - DB credentials
    - Speech keys
    - OpenAI keys
    - storage keys

Audit

Log:

- login
- upload
- transcription start/end
- generation start/end
- each regeneration
- each edit/finalization
- each download/delete

Data minimization

Since you said patient details won't really be stored in MVP:

- avoid DOB, diagnosis, insurance, etc.
- store only minimum patient reference info needed for therapist workflow

Retention

Use **Blob lifecycle policies** for 1-year retention workflows, with policy-based deletion/transition support in storage. (Microsoft Learn)

---

7. Recommended design choices I'd lock in now

Use these

- Flutter for iPhone + web
- FastAPI
- Azure Container Apps
- Azure Database for PostgreSQL
- Azure Blob Storage
- Azure Service Bus
- Azure AI Speech batch transcription
- Azure OpenAI
- Azure Key Vault
- Entra External ID

Avoid these in MVP

- AKS
- microservices split into 10 deployables
- diarization
- patient self-service portal
- EHR integrations
- prior-note retrieval
- multi-practice SaaS complexity
- live streaming transcription

That keeps the MVP within your budget and pilot scope.

---

8. Architecture recommendation summary

My **best-fit MVP architecture** is:

- **Frontend:** Flutter mobile + Flutter web
- **Auth:** Entra External ID
- **API:** FastAPI in Azure Container Apps
- **Async orchestration:** Service Bus + worker containers
- **Storage:** Blob Storage for files, PostgreSQL for metadata/audit
- **Transcription:** Azure AI Speech batch transcription
- **LLM:** Azure OpenAI
- **Docs:** server-side DOCX generation
- **Security:** Key Vault, RBAC, audit logs, 1-year retention

The one place I'd push back on your current answers is **speaker diarization**: for therapist-only post-session recordings, I'd leave it out of MVP.

Next, I'd suggest I turn this into a **detailed Azure service design** with:

- exact Azure resources
- deployment topology
- environment layout
- request flow
- session/job state machine
- and a simple architecture diagram you can use in a pitch or design doc.